

```

1 struct ThreadContext {
2     std::vector<void*> pending_reclaims;
3     std::atomic<uint64_t> counter;
4     ThreadContext *next;
5
6     ThreadContext(MemManager m) {
7         while (true) {
8             next = m.head;
9             if (m.head.compare_exchange_strong(next, this))
10                 break;
11         }
12     }
13 }
14 struct MemManager {
15     static thread_local ThreadContext *self = nullptr;
16     std::atomic<ThreadContext*> head;
17     ...
18 }
19 MemManager::register_thread(int num) { self = new ThreadContext(this); }
20 MemManager::unregister_thread() { /* no-op */ }
21 MemManager::op_begin() { self->counter++; }
22 void MemManager::sched_for_reclaim(void* ptr) { self->pending_reclaims.push_back(ptr); }
23 bool MemManager::try_reserve(void* ptr) { return false; }
24 void MemManager::unreserve(void* ptr) { }
25
26 void MemManager::op_end() {
27     self->counter++;
28     if (pending_reclaims.count() == 0)
29         return;
30     wait_until_unreserved()
31     for (auto p : pending_reclaims)
32         free(p);
33 }
34 void MemManager::wait_until_unreserved() {
35     ThreadContext* curr = head;
36     while (curr) {
37         uint64_t val = curr->counter;
38         if (odd(val))
39             do {
40                 wait();
41             } while (curr->counter.read() == val)
42             curr = curr->next;
43     }
44 }
```

FIGURE 19.8 Epoch-based reclamation.