

The Art of Multiprocessor Programming
Solutions to Exercises
Chapter 8

July 14, 2009

```

1 public class BadCLHLock implements Lock {
2     // most recent lock holder
3     AtomicReference<Qnode> tail;
4     // thread-local variable
5     ThreadLocal<Qnode> myNode;
6     public void lock() {
7         Qnode qnode = myNode.get();
8         qnode.locked = true;           // I'm not done
9         // Make me the new tail, and find my predecessor
10        Qnode pred = tail.getAndSet(qnode);
11        // spin while predecessor holds lock
12        while (pred.locked) {}
13    }
14    public void unlock() {
15        // reuse my node next time
16        myNode.get().locked = false;
17    }
18    static class Qnode { // Queue node inner class
19        public boolean locked = false;
20    }
21 }
```

Figure 1: An incorrect attempt to implement a CLHLock.

Exercise 85. Fig. 1 shows an alternative implementation of CLHLock in which a thread reuses its own node instead of its predecessor node. Explain how this implementation can go wrong.

Solution Suppose A acquires and releases the lock. At this point, the same node is referenced by `tail` and `myNode`. If A tries to reacquire the lock, it sets the node's `locked` field to `true` (Line 8), swaps that node with itself (Line 10), implicitly making the node its own predecessor. It then deadlocks waiting for that node to become unlocked.