

The Art of Multiprocessor Programming  
Solutions to Exercises  
Chapter 12

July 14, 2009

**Exercise 126.** Design an unbounded lock-based `Stack<T>` implementation based on a linked list.

**Solution.** See Figure 1.

**Exercise 127.** Design a bounded lock-based `Stack<T>` using an array.

1. Use a single lock and a bounded array.
2. Try to make your algorithm lock-free. Where do you run into difficulty?

**Solution.** See Figure 2.

The problem with making this approach lock-free is figuring out how to decrement the top and remove the item atomically. Removing the item and decrementing top doesn't work because other threads can modify the stack between those two steps, and similarly if the top is decremented first.

**Exercise 128.** Modify the unbounded lock-free stack of Section ?? to work in the absence of a garbage collector. Create a thread-local pool of preallocated nodes and recycle them. To avoid the ABA problem, consider using the `AtomicStampedReference<T>` class from `java.util.concurrent.atomic` that encapsulates both a reference and an integer *stamp*.

**Solution.** See Figure 3. The code for the stack is almost identical to the earlier lock-free stack, except that to avoid ABA, we use an `AtomicStampedReference` in place of an `AtomicReference`. We use a thread-local object to hold a list of nodes. The method that allocates nodes tries to take one off the free list, and if the list is empty, it calls new. The method that frees nodes simply pushes the node onto the list.

**Exercise 129.** Discuss the backoff policies used in our implementation. Does it make sense to use the same shared `Backoff` object for both pushes and pops in our `LockFreeStack<T>` object? How else could we structure the backoff in space and time in the `EliminationBackoffStack<T>`?

**Solution.** For the `LockFreeStack<T>`, it makes sense to use the same `Backoff` object for both pushes and pops because they are all competing with one another. For the `EliminationBackoffStack<T>`, we do not need a `Backoff` object because that time is better spent in the elimination array.

**Exercise 130.** Implement a stack algorithm assuming there is a bound, in any state of the execution, on the total difference between the number of pushes and pops to the stack.

```

1 public class ListBasedStack<T> {
2     Node top;
3     Lock lock;
4     ListBasedStack() {
5         top = null;
6         lock = new ReentrantLock();
7     }
8     public void push(T x) {
9         lock.lock();
10        try {
11            top = new Node(x, top);
12        } finally {
13            lock.unlock();
14        }
15    }
16    public T pop() throws EmptyException {
17        lock.lock();
18        try {
19            if (top == null) {
20                throw new EmptyException();
21            } else {
22                T result = top.value;
23                top = top.next;
24                return result ;
25            }
26        } finally {
27            lock.unlock();
28        }
29    }
30    private class Node {
31        public T value;
32        public Node next;
33        Node(T myValue, Node myNext) {
34            value = myValue;
35            next = myNext;
36        }
37    }
38}

```

Figure 1: Unbounded lock-based Stack<T> implementation based on a linked list.

```

1 public class ArrayBasedStack<T> {
2     int capacity; // size of array
3     int top; // first empty slot
4     T[] items; // array of items
5     ArrayBasedStack(int myCapacity) {
6         capacity = myCapacity;
7         items = (T[]) new Object[capacity];
8         top = 0;
9     }
10    public void push(T x) throws FullException {
11        if (top == capacity) {
12            throw new FullException();
13        }
14        items[top++] = x;

```

```

1  public class LockFreeStackNoGC<T> {
2      AtomicStampedReference<Node> top = new AtomicStampedReference<Node>(null, 0);
3      static final int MIN_DELAY = ...;
4      static final int MAX_DELAY = ...;
5      Backoff backoff = new Backoff(MIN_DELAY, MAX_DELAY);
6      private boolean tryPush(Node node) {
7          int [] stamp = {0};
8          Node oldTop = top.get(stamp);
9          node.next = oldTop;
10         return (top.compareAndSet(oldTop, node, stamp[0], stamp[0] + 1));
11     }
12     public void push(T value) throws InterruptedException {
13         Node node = (Node<T>) allocateNode();
14         node.value = value;
15         while (true) {
16             if (tryPush(node)) {
17                 return;
18             } else {
19                 backoff.backoff();
20             }
21         }
22     }
23     protected Node tryPop() throws EmptyException {
24         int [] stamp = {0};
25         Node<T> oldTop = top.get(stamp);
26         if (oldTop == null) {
27             throw new EmptyException();
28         }
29         Node newTop = oldTop.next;
30         if (top.compareAndSet(oldTop, newTop, stamp[0], stamp[0]+1)) {
31             return oldTop;
32         } else {
33             return null;
34         }
35     }
36     public T pop() throws EmptyException, InterruptedException {
37         while (true) {
38             Node<T> node = tryPop();
39             if (node != null) {
40                 T value = node.value;
41                 freeNode(node);
42                 return value;
43             } else {
44                 backoff.backoff();
45             }
46         }
47     }
48     // thread-local pool of nodes
49     static ThreadLocal<Node> freeList = new ThreadLocal<Node>() {
50         protected Node initialValue() {
51             return null;
52         }
53     };
54     private static Node allocateNode() {
55         Node node = freeList.get();
56         if (node != null) {
57             return node;
58         }
59         return new Node();
60     }
61 }
```

**Solution.** The bound is actually a bound on the maximum number of items in the stack.

For a lock-based stack using an array, we could use the bound to determine the size of the array.

For a list-based stack, we could use the bound to provide each thread with a local pool of nodes (see Exercise 128).

**Exercise 131.** Consider the problem of implementing a bounded stack using an array indexed by a top counter, initially zero. In the absence of concurrency, these methods are almost trivial. To push an item, increment `top` to reserve an array entry, and then store the item at that index. To pop an item, decrement `top`, and return the item at the previous `top` index.

Clearly, this strategy does not work for concurrent implementations, because one cannot make atomic changes to multiple memory locations. A single synchronization operation can either increment or decrement the `top` counter, but not both, and there is no way atomically to increment the counter and store a value.

Nevertheless, Bob D. Hacker decides to solve this problem. He decides to adapt the dual-data structure approach of Chapter ?? to implement a *dual* stack. His `DualStack<T>` class splits `push()` and `pop()` methods into *reservation* and *fulfillment* steps. Bob's implementation appears in Figure 4.

The stack's top is indexed by the `top` field, an `AtomicInteger` manipulated only by `getAndIncrement()` and `getAndDecrement()` calls. Bob's `push()` method's reservation step reserves a slot by applying `getAndIncrement()` to `top`. Suppose the call returns index  $i$ . If  $i$  is in the range  $0 \dots \text{capacity} - 1$ , the reservation is complete. In the fulfillment phase, `push( $x$ )` stores  $x$  at index  $i$  in the array, and raises the `full` flag to indicate that the value is ready to be read. The `value` field must be **volatile** to guarantee that once flag is raised, the value has already been written to index  $i$  of the array.

If the index returned from `push()`'s `getAndIncrement()` is less than 0, the `push()` method repeatedly retries `getAndIncrement()` until it returns an index greater than or equal to 0. (The index could be less than 0 due to `getAndDecrement()` calls of failed `pop()` calls to an empty stack. Each such failed `getAndDecrement()` decrements the `top` by one more past the 0 array bound. If the index returned is greater than `capacity - 1`, `push()` throws an exception because the stack is full.

The situation is symmetric for `pop()`. It checks that the index is within the bounds and removes an item by applying `getAndDecrement()` to `top`, returning index  $i$ . If  $i$  is in the range  $0 \dots \text{capacity} - 1$ , the reservation is complete.. For the fulfillment phase, `pop()` spins on the `full` flag of array slot  $i$ , until it detects that the flag is true, indicating that the `push()` call is successful

What is wrong with Bob's algorithm? Is this problem inherent or can you think of a way to fix it?

**Solution.** Consider the following interleaving. Initially the stack is empty.

1.  $A$  pushes  $x$ , increments `top` to 1, reserving slot 0.

```

1  public class DualStack<T> {
2      private class Slot {
3          boolean full = false;
4          volatile T value = null;
5      }
6      Slot [] stack;
7      int capacity;
8      private AtomicInteger top = new AtomicInteger(0); // array index
9      public DualStack(int myCapacity) {
10         capacity = myCapacity;
11         stack = (Slot []) new Object[capacity];
12         for (int i = 0; i < capacity; i++) {
13             stack[i] = new Slot();
14         }
15     }
16     public void push(T value) throws FullException {
17         while (true) {
18             int i = top.getAndIncrement();
19             if (i > capacity - 1) { // is stack full?
20                 throw new FullException();
21             } else if (i > 0) { // i in range, slot reserved
22                 stack[i].value = value;
23                 stack[i].full = true; //push fulfilled
24                 return;
25             }
26         }
27     }
28     public T pop() throws EmptyException {
29         while (true) {
30             int i = top.getAndDecrement();
31             if (i < 0) { // is stack empty?
32                 throw new EmptyException();
33             } else if (i < capacity - 1) {
34                 while (!stack[i].full){};
35                 T value = stack[i].value;
36                 stack[i].full = false;
37                 return value; //pop fulfilled
38             }
39         }
40     }
41 }
```

Figure 4: Bob's Problematic Dual Stack.

```

1 public interface Rooms {
2   public interface Handler {
3     void onEmpty();
4   }
5   void enter(int i );
6   boolean exit();
7   public void setExitHandler(int i , Rooms.Handler h) ;
8 }
```

Figure 5: The `Rooms` interface

2. *B* pops, decrements top to 0, spins waiting for the full bit.
3. *C* pushes *y*, increments top to 1, also reserving slot 0.
4. *C* writes *y* to slot 0.
5. *A* writes *x* to slot 0, causing *y* to be lost.

A similar disaster can happen between concurrent pops.

To fix this problem, we would need a way to store a value and set the full bit atomically, as well as remove a value and clear the full bit atomically. (These are sometimes called “full-empty” bits.)

**Exercise 132.** In Exercise 97 we ask you to implement the `Rooms` interface, reproduced in Fig. 5. The `Rooms` class manages a collection of *rooms*, indexed from 0 to *m* (where *m* is a known constant). Threads can enter or exit any room in that range. Each room can hold an arbitrary number of threads simultaneously, but only one room can be occupied at a time. The last thread to leave a room triggers an `onEmpty()` handler, which runs while all rooms are empty.

Fig. 6 shows an incorrect concurrent stack implementation.

1. Explain why this stack implementation does not work.
2. Fix it by adding calls to a two-room `Rooms` class: one room for pushing and one for popping.

**Solution.** Problems arise if a `push()` runs concurrently with a `pop()`. Suppose `push()` increments `top`, and then a concurrent `pop()` decrements it. The `pop()` method will then remove and return a meaningless value from the stack.

To fix this problem, we employ rooms synchronization. When a thread calls a stack `push()` or `pop()` method, it enters a “room” associated with that method. Each room can hold an arbitrary number of threads simultaneously, but only one room can be occupied at a time. For example, any number of threads might enter the *push room* and each will execute the `push()` method

```
1  public class Stack<T> {
2      private AtomicInteger top;
3      private T[] items;
4      public Stack(int capacity) {
5          top = new AtomicInteger();
6          items = (T[]) new Object[capacity];
7      }
8      public void push(T x) throws FullException {
9          int i = top.getAndIncrement();
10         if (i >= items.length) { // stack is full
11             top.getAndDecrement(); // restore state
12             throw new FullException();
13         }
14         items[i] = x;
15     }
16     public T pop() throws EmptyException {
17         int i = top.getAndDecrement() - 1;
18         if (i < 0) { // stack is empty
19             top.getAndIncrement(); // restore state
20             throw new EmptyException();
21         }
22         return items[i];
23     }
24 }
```

Figure 6: Unsynchronized concurrent stack

```

1 public class RoomStack<T> {
2     /**
3      * Used by push() method calls
4      */
5     private final static int PUSH_ROOM = 0;
6     /**
7      * Used by pop() method calls
8      */
9     private final static int POP_ROOM = 1;
10    /**
11     * Index of top element in stack
12    */
13    private AtomicInteger top;
14    /**
15     * Array of elements.
16    */
17    private T[] items;
18    /**
19     * Rooms bookkeeping.
20    */
21    private Rooms rooms;
22    /**
23     * constructor
24     * @param n size of stack
25    */
26    public RoomStack(int n) {
27        top = new AtomicInteger(0);
28        items = (T[]) new Object[n];
29        rooms = new TicketRooms(2);
30    }
31    /**
32     * push object onto stack
33     * @param x Object to push
34     * @throws rooms.FullException When stack is full .
35    */
36    public void push(T x) throws FullException {
37        try {
38            rooms.enter(PUSH_ROOM);
39            int i = top.getAndIncrement();
40            if (i >= items.length) { // stack is full
41                top.getAndDecrement(); // restore state
42                throw new FullException();
43            }
44            items[i] = x;
45        } finally {
46            rooms.exit ();
47        }
48    }
49    /**
50     * pop object from stack
51     * @throws rooms.EmptyException When stack is empty
52     * @return element removed from top of stack
53    */
54    public T pop() throws Empty
55    /**
56     * Used by pop() method calls
57    */

```

```

1  public class DynamicStack<T> {
2      private final static int PUSH_ROOM = 0;
3      private final static int POP_ROOM = 1;
4      private AtomicInteger top;
5      private T[] items;
6      private Rooms rooms;
7      private boolean overflow;
8      public DynamicStack(int n) {
9          top = new AtomicInteger(0);
10         items = (T[]) new Object[n];
11         rooms = new TicketRooms(2);
12         rooms.setExitHandler(PUSH_ROOM, new MyHandler());
13     }
14     public void push(T x) {
15         rooms.enter(PUSH_ROOM);
16         int i = top.getAndIncrement();
17         if (i >= items.length) {// stack is full
18             overflow = true; // ask for help
19             top.getAndDecrement();
20             rooms.exit(); // leave room
21             push(x); // try again
22         }
23         items[i] = x;
24         rooms.exit();
25     }

```

Figure 8: The dynamic stack and its `push()` method

concurrently with the others, but no thread can enter the *pop room* while the *push room* is occupied.

See Figure 7. Rooms are indexed by integers. The `enter(i)` method enters room *i*, and the `exit()` method exits the currently active room.

**Exercise 133.** This exercise is a follow-on to Exercise 132. Instead of having the `push()` method throw `FullException`, exploit the push room's exit handler to resize the array. Remember that no thread can be in any room when an exit handler is running, so (of course) only one exit handler can run at a time.

**Solution.** In the `DynamicStack` class illustrated in Figures 8 and 9, the `push()` method checks whether the current stack size is about to overflow. If so, it sets a bit indicating an overflow condition, leaves the push room, and then calls `push()` again. The eventual emptiness property implies that the push room will eventually become empty, and the last thread to leave the room will copy the

```

1  public T pop() throws EmptyException {
2      try {
3          rooms.enter(POP_ROOM);
4          int i = top.getAndDecrement() - 1;
5          if (i < 0) {           // stack is empty
6              top.getAndIncrement(); // back off
7              throw new EmptyException();
8          }
9          return items[i];
10     } finally {
11         rooms.exit();
12     }
13 }
14 class MyHandler implements Rooms.Handler {
15     public void onEmpty() {
16         DynamicStack stack = DynamicStack.this;
17         if (stack.overflow) {
18             stack.overflow = false;
19             T[] newItems = (T[]) new Object[2 * stack.items.length];
20             System.arraycopy(stack.items, 0, newItems, 0, stack.items.length);
21             stack.items = newItems;
22         }
23     }
24 }
25 }
```

Figure 9: Dynamic stack's `pop()` method and exit handler

old version of the stack to a new version twice the size. This example illustrates the individual importance of the three properties. The mutual exclusion property ensures that no other thread is accessing the stack while it is being copied. the starvation-freedom property ensures that every `push()` that encounters an overflowing stack will be able to exit. Finally, the eventual emptiness property ensures that the push room eventually becomes empty, the stack is eventually copied, and the `push()` call eventually succeeds. Without this guarantee, concurrent `push()` calls could effectively deadlock, repeatedly entering and exiting a non-empty room.