

The Art of Multiprocessor Programming
Solutions to Exercises
Chapter 14

July 14, 2009

Exercise 163. Recall that a skiplist is a *probabilistic* data structure. Although the expected performance of a `contains()` call is $O(\log n)$, where n is the number of items in the list, the worst-case performance could be $O(n)$. Draw a picture of an 8-element skiplist with worst-case performance, and explain how it got that way.

Solution See Fig. 1. Here, when value i was inserted, it randomly chose height $i + 1$.

Exercise 164. You are given a skiplist with probability p and `MAX_LEVEL` M . If the list contains N nodes, what is the expected number of nodes at each level from 0 to $M - 1$?

Solution The expected number of nodes at level i or higher is $p^i \gg N$. The expected number at level i is $f(p^i - p^{i+1}) \cdot N$, for $i < M - 1$.

Exercise 165. Modify the `LazySkipList` class so `find()` starts at the level of the highest node currently in the structure, instead of the highest level possible (`MAX_LEVEL`).

Solution See the `LazySkipListLimit` class in the code folder. We keep the highest actual level in an `AtomicInteger` field, which we update whenever we create a node higher than the previous maximum. In this solution, we do not try to decrease the max if the node at the highest level is removed.

Exercise 166. Modify the `LazySkipList` to support multiple items with the same key.

Solution Simply change `add()` to return `void`, and delete the code where `add()` checks whether the value is already present.

Exercise 167. Suppose we modify the `LockFreeSkipList` class so that at Line ?? of Fig. ??, `remove()` restarts the main loop instead of returning `false`.

Is the algorithm still correct? Address both safety and liveness issues. That is, what is an unsuccessful `remove()` call's new linearization point, and is the class still lock-free?

Solution The method is still correct, but no longer lock-free. The unmodified method had three possible linearization points, corresponding to finding the item, not finding the item, or finding the item marked. The modification simply eliminates the third. The class is no longer lock-free because a thread that halts after marking an item will block any threads looking for that item.




Figure 1: Worst-case skiplist performance

Exercise 168. Explain how, in the `LockFreeSkipList` class, a node might end up in the list at levels 0 and 2, but not at level 1. Draw pictures.

Solution See Fig. 3. At the start, there are two nodes, 2 and 5, both at level 2.

- Thread *A* inserts node 3 at level 2 between nodes 2 and 5.
- *B* removes 3, marking the references from 2 to 3 at levels 0 and 1, but then pauses.
- *C* inserts 4. It searches at level 2, finds an unmarked link from 3 to 5, so it moves down to level 1. At level 1, it finds a marked reference from 2 to 3, so it physically deletes that node, redirecting it from 2 to 5.

Exercise 169. Modify the `LockFreeSkipList` so that the `find()` method snips out a sequence of marked nodes with a single `compareAndSet()`. Explain why your implementation cannot remove a concurrently inserted unmarked node.

Solution [[[Nir?]]]

Exercise 170. Will the `add()` method of the `LockFreeSkipList` work even if the bottom level is linked and then all other levels are linked in some arbitrary order? Is the same true for the marking of the `next` references in the `remove()` method: the bottom level `next` reference is marked last, but references at all other levels are marked in an arbitrary order?

```

1  boolean contains(T x) {
2      int bottomLevel = 0;
3      int key = x.hashCode();
4      Node<T> pred = head;
5      Node<T> curr = null;
6      for (int level = MAX_LEVEL; level >= bottomLevel; level--) {
7          curr = pred.next[ level ].getReference();
8          while (curr.key < key) {
9              pred = curr;
10             curr = pred.next[ level ].getReference();
11         }
12     }
13     return curr.key == key;
14 }
```

Figure 2: The `LockFreeSkipList` class: an *incorrect* `contains()`.




Figure 3: Solution to Exercise 168

Solution [[[Nir?]]]

Exercise 171. (Hard) Modify the `LazySkipList` so that the list at each level is bidirectional, and allows threads to add and remove items in parallel by traversing from either the `head` or the `tail`.

Solution [[[Nir?]]]

Exercise 172. Fig. 2 shows a buggy `contains()` method for the `LockFreeSkipList` class. Give a scenario where this method returns a wrong answer. Hint: the reason this method is wrong is that it takes into account keys of nodes that have been removed.

Solution

Solution 172 Initially the list has height 2, and has the values 1 and 4 in both levels.

1. A call to `add(2)` links a node with key 2 in the bottom layer, and stalls (Fig. 4).
2. A call to `add(2)` links a node with key 2 in the bottom layer, and stalls (Fig. 5).
3. A call to `remove(2)` marks 2's node at the bottom layer (Fig. 6).
4. The `add(2)` call continues. It fails to link 2's node in the 2nd layer, so it calls `find()` again, which now returns 1 and 3 as predecessor and successor. 2 is successfully spliced in between them (Fig. 7).

Now a `contains(3)` call finds 2 and 3 as the predecessor and successor in the second layer, goes down to the bottom layer, and traverses 2's `next` reference to 4, and returns that 3 is not in the list. This execution is not linearizable, because the call to `add(3)` has returned.



Figure 4: Counterexample: Step 1




Figure 5: Counterexample: Step 2




Figure 6: Counterexample: Step 3




Figure 7: Counterexample: Step 4